

Hi and welcome to the Covera NLP exercise!!

Data Science Exercise

Hello! On behalf of Covera Health, we are excited that you are considering joining our team!

Part of your interview process consists in solving a data science exercise.

Please make sure you have gone through the setup process as defined in the README.md.

We really hope that you will enjoy working on this exercise with us; it's a small glimpse into what we do here on a day-to-day basis! Have fun! :)

Background/Context

Radiology has been transformative in patient care. Being able to identify pathologies for diagnoses through advanced technologies such as MRI and CT has helped countless lives. However, Radiology is a very complicated discipline, and even the most qualified radiologists make mistakes from time to time while interpreting images and describing their findings in diagnostic reports. For example, a recent study published in the Spine Journal demonstrated error rates of lumbar MRIs to be as high as 43%. These errors in diagnosis have significant impact to patient care, as all of the downstream treatment decisions are predicated on the initial diagnosis.

Covera Health is focused on improving patient outcomes by estimating radiology provider error rates. We rely on expert reviews of selected medical diagnostic reports to inform our error rate estimates. Our task is to analyse diagnostic reports automatically to augment the expert reviews.

A first step in what we do is to extract and identify the diseases mentioned in a biomedical test. For the purposes of this exercise In this notebook, we'll be using a dataset from the [NCBI disease database](https://www.ncbi.nlm.nih.gov/CBBresearch/Dogan/DISEASE/) (<https://www.ncbi.nlm.nih.gov/CBBresearch/Dogan/DISEASE/>), which is publicly available and contains 6,293 English sentences from 793 PubMed abstracts. Each sentence has been annotated by the named diseases they contain. Our goal to is develop a model which can automatically extract named diseases from the text.

1. Let's load the data and give it a browse

```
In [1]: import pandas as pd

data = pd.read_csv('data/ncbi_disease_dataset.csv',encoding='latin1')

data.fillna(method="ffill",inplace=True)
print('Total Number of Sentences: ',data['Sentence #'].unique().shape[0])
```

Total Number of Sentences: 6293

```
In [2]: data.head(10)
```

Out[2]:

	Word	Tag	Sentence #
0	Identification	O	Sentence :0
1	of	O	Sentence :0
2	APC2	O	Sentence :0
3	,	O	Sentence :0
4	a	O	Sentence :0
5	homologue	O	Sentence :0
6	of	O	Sentence :0
7	the	O	Sentence :0
8	adenomatous	B-Disease	Sentence :0
9	polyposis	I-Disease	Sentence :0

```
In [3]: data['Tag'].value_counts()/data['Tag'].value_counts().sum()
```

```
Out[3]: O          0.917684
I-Disease  0.044972
B-Disease  0.037344
Name: Tag, dtype: float64
```

Looks like we've got a very sparse dataset, with 'O' (outside of entity) being the vast majority of tags, which makes sense. A lot of language is grammar, pronouns, broader ideas, etc. and not named entities.

What are the different entities?

```
In [4]: all_tags = data['Tag'].unique().tolist()
all_tags.remove('O')
all_tags.sort(key = lambda x: x.split('-')[1])
for tag in all_tags:
    print('_____')
    print(tag)
    print('_____')
    print(data[data['Tag']==tag]['Word'].head(20))
```

B-Disease

8	adenomatous
15	adenomatous
61	colon
173	colon
197	cancer
208	HNPCC
221	colorectal
244	hereditary
251	HNPCC
308	colorectal
332	HNPCC
354	colorectal
367	colorectal
476	HNPCC
487	colorectal
534	cancers
537	colorectal
553	colorectal
603	endometrial
615	premenopausal

Name: Word, dtype: object

I-Disease

9	polyposis
10	coli
11	tumour
16	polyposis
17	coli
18	(
19	APC
20)
21	tumour
62	carcinoma
174	carcinoma
222	cancer
245	non
246	-
247	polyposis
248	cancer
249	syndrome
309	cancer
368	cancer
488	,

Name: Word, dtype: object

The NER tags follow standard tagging convention defined [here](https://en.wikipedia.org/wiki/Inside%E2%80%93outside%E2%80%93beginning_(tagging)) ([https://en.wikipedia.org/wiki/Inside%E2%80%93outside%E2%80%93beginning_\(tagging\)](https://en.wikipedia.org/wiki/Inside%E2%80%93outside%E2%80%93beginning_(tagging))) and B-disease is the beginning of a disease tag and an I-Disease in inside a disease tag. For eg: **colorectal cancer** is comprised of two separate tags as follows: colorectal: **B-Disease** and cancer: **I-Disease** .

```
In [5]: data['Tag'].value_counts()/data['Tag'].value_counts().sum()
```

```
Out[5]: O                0.917684  
I-Disease    0.044972  
B-Disease    0.037344  
Name: Tag, dtype: float64
```

```
In [6]: data.describe(  
)
```

```
Out[6]:
```

	Word	Tag	Sentence #
count	184447	184447	184447
unique	10798	3	6293
top	.	O	Sentence :981
freq	8361	169264	204

2. Preparing text for model inputs

Question 1: How should you prepare the data for modeling?

```
In [7]: # since we are doing sentence level, we will make the each value in X a  
        # list of the tokens for each sentence  
        # and for y we will make each value a list of the tags for the tokens in  
        # the corresponding sentence  
  
        # we will also lower each token  
  
X = data.groupby('Sentence #').Word.apply(lambda x: [s.lower() for s in  
x])  
y = data.groupby('Sentence #').Tag.apply(list)
```

```
In [8]: X.head()
```

```
Out[8]: Sentence #  
Sentence :0      [identification, of, apc2, ,, a, homologue, of...  
Sentence :1      [., the, adenomatous, polyposis, coli, (, apc,...  
Sentence :10     [., a, common, msh2, mutation, in, english, an...  
Sentence :100    [., we, report, that, heterozygosity, for, a, ...  
Sentence :1000   [., eleven, of, the, 16, homozygous, subjects,...  
Name: Word, dtype: object
```

```
In [9]: y.head()
```

```
Out[9]: Sentence #
Sentence :0      [0, 0, 0, 0, 0, 0, 0, 0, 0, B-Disease, I-Disease,...
Sentence :1      [0, 0, B-Disease, I-Disease, I-Disease, I-Dise...
Sentence :10     [0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, B-Disease, 0, 0...
Sentence :100    [0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, ...
Sentence :1000   [0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, B-Dise...
Name: Tag, dtype: object
```

```
In [10]: #Remove cell
#We're going to be creating a sequence based model with batching, which
means we need to define a maximum sequence length
max_len = 75
```

```
In [11]: data.groupby('Sentence #').Word.count().hist()
```

```
Out[11]: <matplotlib.axes._subplots.AxesSubplot at 0x7febf06150d0>
```

```
In [12]: # To make the word2idx dictionary we first get all of the words and map
the to lower strings
# then we sort by their frequency (in case we want to drop infrequent to
kens and reduce our vocab size)
# This will give us a dataframe with the word as the index, and the coun
ts as the column
# In order to give it a unique idx, we create a new column where the val
ue is just the order of the word
# now we can convert this to a dictionary, where the key is a word and t
he value is its unique index

word_counts = data.Word.dropna().map(str.lower).value_counts().rename('c
ounts').to_frame()
word_counts['index_num'] = range(1, word_counts.shape[0]+1)
word2idx = word_counts['index_num'].to_dict()

# we'll set a special idx for tokens that were not part of the original
vocab
word2idx['UNK'] = 0
# we'll also set a special idx to pad our sentence vectors
word2idx['PAD'] = -1
```

```
In [16]: tag_counts = data.Tag.dropna().value_counts().rename('counts').to_frame
()
tag_counts['index_num'] = range(0, tag_counts.shape[0])
tag2idx = tag_counts['index_num'].to_dict()

tag2idx['PAD'] = -1

tag2idx
```

```
Out[16]: {'O': 0, 'I-Disease': 1, 'B-Disease': 2, 'PAD': -1}
```

Question 2: Why do we need a max length? We set the maximum sequence length to 75 words, do you agree with this choice, why or why not?

Assuming we're translating words and labels to indices write some code to create a hash table that can map our strings to numerical values.

```
In [17]: # word2idx =  
# tag2idx =
```

```
In [18]: from keras.preprocessing.sequence import pad_sequences  
  
import numpy as np  
X_word = [[word2idx[w] for w in s] for s in X.tolist()]  
X_word = pad_sequences(maxlen=max_len, sequences=X_word, value=word2idx[  
"PAD"], padding='post', truncating='post')
```

```
In [ ]:
```

We'll do the same thing for our TAG values which are the labels for the named entities

```
In [19]: y = [[tag2idx[w] for w in s] for s in y]  
y = pad_sequences(maxlen=max_len, sequences=y, value=tag2idx["PAD"], padding='post', truncating='post')
```

Bonus question: let's assume there are lots of misspellings in our training data, and we have no way to correct them before they get to the model. How would you deal with this if at all?

```
In [ ]: # As opposed to something like a dictionary-based method where you might  
# be able  
# to return a result for an exact phrase match,  
# Instead statistical named entity recognition is more robust to misspell  
# ings b/c the model  
# learns to predict the tag not just on the token itself but also on the  
# neighboring tokens as well
```

We've got the data into a format that can be fed into our model. Now let's decide how to split our data to train, validate and test our model. Write some code to split the data.

```
In [ ]: from sklearn
def create_splits(X, Y, ):
    """creates function that takes in X, Y"""
    pass

X_train =
X_test =
X_val =
y_train =
y_test =
y_val =
```

Question 3: Now that we've got our data prepared, we're going to start building a model. Are there any additional steps we should take before feeding our data to a model?

```
In [ ]: # if we are utilizing the tokens as is, I would utilize
```

Question 4: How do we ensure that our model generalizes and that our metrics are representative of real world data?

Question 5: How would you select the best model for the problem?

- What are some considerations when selecting the best model?
- What are some things you can do to ensure that the model you've selected is indeed the best one?

Build a Model

```
In [ ]: def create_model():
    model = None
    return model
```

Question 6: Explain whether we should use the following code, and if so, why?


```

In [ ]: from keras.callbacks import EarlyStopping, ModelCheckpoint
        from keras.models import load_model

        model = create_model()

        callbacks = [
            EarlyStopping(
                monitor='val_loss',
                patience=3
            ),
            ModelCheckpoint(
                filepath='models/ner_checkpoint_test.h5',
                monitor='val_loss',
                save_best_only=True
            )
        ]

        history = model.fit(X_train,
                            y_train,
                            validation_split=0.1, #change to validation data
                            callbacks = callbacks,
                            verbose=1
        )

```

3. Validation

Now let's assume we've got predictions out of our model. We'll load some dummy predictions to evaluate the performance of our hypothetical model.

```

In [15]: #load in prepared predictions and evaluate how well we predicted
        predictions = pd.read_csv('predictions/ncbi_disease_predictions.csv')
        tag2idx = dict()
        for idx, i in enumerate(sorted(predictions['Tag'].unique().tolist())):
            tag2idx[i] = idx
        predictions['true_index'] = predictions['Tag'].map(tag2idx)
        predictions['predicted_index'] = predictions['predicted'].map(tag2idx)

```

```
In [16]: predictions.head()
```

```
Out[16]:
```

	Unnamed: 0	Word	Tag	Sentence #	true_index	predicted_index	predicted
0	0	Clustering	O	Sentence :0	2	2	O
1	1	of	O	Sentence :0	2	2	O
2	2	missense	O	Sentence :0	2	2	O
3	3	mutations	O	Sentence :0	2	2	O
4	4	in	O	Sentence :0	2	2	O

And we'll use a confusion matrix to see how we did on each prediction.

```

In [17]: import matplotlib.pyplot as plt
import numpy as np
import itertools

def plot_confusion_matrix(cm,
                           classes,
                           normalize=False,
                           title='Confusion matrix',
                           save_file = 'confusion_matrix',
                           cmap=plt.cm.Blues):
    """
    This function prints and plots the confusion matrix.
    Normalization can be applied by setting `normalize=True`.
    --inputs--
    cm - the confusion matrix produced by sklearn's confusion_matrix func
tion
    classes - the labels for the classes in the confusion matrix

    --outputs--
    matplotlib figure displaying the confusion matrix, and the saved png
of the figure

    """
    if normalize:
        cm = cm.astype('float') / cm.sum(axis=1)[:, np.newaxis]
        print("Normalized confusion matrix")
    else:
        print('Confusion matrix, without normalization')

    #print(cm)
    plt.figure(figsize=(8,8))
    plt.imshow(cm, interpolation='nearest', cmap=cmap)
    plt.title(title)
    plt.colorbar()
    tick_marks = np.arange(len(classes))
    plt.xticks(tick_marks, classes, rotation=90)
    plt.yticks(tick_marks, classes)

    fmt = '.2f' if normalize else 'd'
    thresh = cm.max() / 2.
    for i, j in itertools.product(range(cm.shape[0]), range(cm.shape[1]
)):
        plt.text(j, i, format(cm[i, j], fmt),
                 horizontalalignment="center",
                 color="white" if cm[i, j] > thresh else "black")

    plt.ylabel('True label')
    plt.xlabel('Predicted label')
    plt.tight_layout()
    plt.savefig('figures/{}.png'.format(save_file))
    plt.show()

```

```
In [18]: from sklearn.metrics import confusion_matrix
def switch_keys_values_dict(d):
    switch_dict = dict()
    for t in tag2idx.keys():
        switch_dict[d[t]] = t
    return switch_dict


idx2tag = switch_keys_values_dict(tag2idx)

labels = list()
for i in range(0, len(idx2tag)):
    labels.append(idx2tag[i])

test_labels = predictions['true_index'].values
pred_labels = predictions['predicted_index'].values

cnf_matrix = confusion_matrix(test_labels, pred_labels)
cnf_title = 'NER Prediction Module'
save_all = 'confusion_matrix_all'
plot_confusion_matrix(
    cnf_matrix,
    classes=labels,
    normalize=True,
    title=cnf_title,
    save_file=save_all
)
```

Normalized confusion matrix



Question 7: Can you think of some single metrics that summarize the performance of our model? What are the tradeoffs we make when we optimize for certain metrics? Can you think of pros and cons of using different metrics?

Question 8: We are seeing some strange results where words that are not diseases are being classified as diseases. The clinical members on the team who are experts in medical language are happy to help you troubleshoot. What would you do to best use their time?

In []:

4. Named Entity Prediction Module

Now that you've built a model, our engineering team needs to use it on raw text to generate predictions. With pseudo code, explain how you would make this model available to teams outside of AI/Data Science

In []:

In []:

In []:

Question 9: What would you do if you had 10x the data? What would you do if you had 1/10 the data?

Question 10: what happens if your model starts doing poorly? How would you understand that your model is doing poorly? How do you measure data drift without labels for new data?

Bonus questions

Bonus question 1: What if instead of predicting the disease Tag we had to predict the disease directly?

Bonus question 2: How would you adapt your model from PubMed abstracts to extracting diseases from clinical notes?

Bonus question 3: Imagine you had access to all the data in the world and an unlimited annotation budget can you design your ideal dataset?

Now that you've built a model, our engineering team needs to use it on raw text to generate predictions. With pseudo code, explain how you would make this model available to teams outside of AI/Data Science

In []:

In []:

In []:

Question 9: What would you do if you had 10x the data? What would you do if you had 1/10 the data?

Question 10: what happens if your model starts doing poorly? How would you understand that your model is doing poorly? How do you measure data drift without labels for new data?

Bonus questions

Bonus question 1: What if instead of predicting the disease Tag we had to predict the disease directly?

Bonus question 2: How would you adapt your model from PubMed abstracts to extracting diseases from clinical notes?

Bonus question 3: Imagine you had access to all the data in the world and an unlimited annotation budget can you design your ideal dataset?

In []: